

MODERN OPERATING SYSTEM

ASSIGNMENT

GURUPRIYAN

20384111

EXPERIMENT - 01

Aim: Write a program to simulate the following non-pre-emptive CPU scheduling algorithms to find turnaround time and waiting time.

(a) FCFS (b) SJF (c) Round Robin (pre-emptive) (d) Priority.

Algorithms:

(a) FCFS (First-Come, First-Served)

1. Order the Processes: List the processes by their arrival time (earliest first).
2. Calculate Waiting and Turnaround Time:
 - First Process: Its waiting time is 0.
 - Next Processes: Add the burst time and waiting time of the previous process to get the waiting time for the current process.
 - Turnaround Time: Add burst time and waiting time for each process.
3. Display Details: Show each process with its ID, burst time, arrival time, waiting time, and turnaround time.

(b) SJF (Shortest Job First)

1. Order the Processes: List the processes by their burst time (shortest first).
2. Calculate Waiting and Turnaround Time: Similar to FCFS.
3. Display Details: Show each process with its ID, burst time, arrival time, waiting time, and turnaround time.

(c) Round Robin (Pre-emptive)

1. Input Processes and Quantum Time.
2. Initialize Times: Set remaining time for each process to its burst time, and waiting and turnaround times to 0.
3. Process Execution: Repeat until all processes are done:
 - Each Process:
 - If its remaining time is \leq quantum and > 0 :
 - Add remaining time to the current time.
 - Set remaining time to 0.
 - Calculate waiting time as current time minus burst time.
 - Turnaround time is the current time.
 - Increment completed count.
 - If remaining time $>$ quantum:

- Add quantum to current time.
 - Subtract quantum from remaining time.
4. Display Details: Show each process with its ID, burst time, arrival time, waiting time, and turnaround time.

(d) Priority

1. Order the Processes: List the processes by their priority (lowest first).
2. Calculate Waiting and Turnaround Time: Similar to FCFS.
3. Display Details: Show each process with its ID, burst time, arrival time, waiting time, and turnaround time.

Source Code:

```
class Process: # Defines a class to represent a process with attributes pid,
arrival_time, burst_time, and priority if provided.
    def __init__(self, pid, arrival_time, burst_time, priority=None):
        self.pid = pid
        self.arrival_time = arrival_time
        self.burst_time = burst_time
        self.priority = priority

def fcfs(processes): # Implements First Come, First Served scheduling
algorithm for given processes.
    processes.sort(key=lambda x: x.arrival_time) # Sort processes by arrival
time
    n = len(processes)
    waiting_time = [0] * n
    turnaround_time = [0] * n
    total_waiting_time = 0
    total_turnaround_time = 0
    for i in range(n):
        if i == 0:
            waiting_time[i] = 0
        else:
            waiting_time[i] = waiting_time[i - 1] + processes[i -
1].burst_time
            turnaround_time[i] = waiting_time[i] + processes[i].burst_time
            total_waiting_time += waiting_time[i]
            total_turnaround_time += turnaround_time[i]
    avg_waiting_time = total_waiting_time / n
    avg_turnaround_time = total_turnaround_time / n
    return avg_waiting_time, avg_turnaround_time

def sjf(processes): # Implements Shortest Job First scheduling algorithm for
given processes.
    processes.sort(key=lambda x: x.burst_time) # Sort processes by burst time
```

```

    return fcfs(processes)

def priority(processes): # Implements Priority scheduling algorithm for given
processes.
    processes.sort(key=lambda x: x.priority, reverse=True) # Sort processes
by priority
    return fcfs(processes)

def round_robin(processes, time_slice): # Implements Round Robin scheduling
algorithm for given processes with a specified time slice.
    n = len(processes)
    burst_remaining = [p.burst_time for p in processes]
    waiting_time = 0
    turnaround_time = 0
    t = 0
    while True:
        done = True
        for i in range(n):
            if burst_remaining[i] > 0:
                done = False
                if burst_remaining[i] > time_slice:
                    t += time_slice
                    burst_remaining[i] -= time_slice
                else:
                    t += burst_remaining[i]
                    waiting_time += t - processes[i].burst_time
                    turnaround_time += t
                    burst_remaining[i] = 0
            if done:
                break
    avg_waiting_time = waiting_time / n
    avg_turnaround_time = turnaround_time / n
    return avg_waiting_time, avg_turnaround_time

if __name__ == "__main__":
    processes = [
        Process(3, 5, 4, 1),
        Process(4, 1, 8, 4),
        Process(3, 2, 3, 1),
        Process(2, 3, 6, 2)
    ]
    print("FCFS:")
    avg_waiting_time, avg_turnaround_time = fcfs(processes)
    print("Average Waiting Time:", avg_waiting_time)
    print("Average Turnaround Time:", avg_turnaround_time)
    print("\nSJF:")
    avg_waiting_time, avg_turnaround_time = sjf(processes)
    print("Average Waiting Time:", avg_waiting_time)

```

```

print("Average Turnaround Time:", avg_turnaround_time)
print("\nPriority:")
avg_waiting_time, avg_turnaround_time = priority(processes)
print("Average Waiting Time:", avg_waiting_time)
print("Average Turnaround Time:", avg_turnaround_time)
print("\nRound Robin:")
avg_waiting_time, avg_turnaround_time = round_robin(processes, 2)
print("Average Waiting Time:", avg_waiting_time)
print("Average Turnaround Time:", avg_turnaround_time)

```

Output:

```

PS C:\Users\D A GURUPRIYAN\Documents\MOS> & "C:/Users/D A GURUPRIYAN/
g_time.py"
FCFS:
Average Waiting Time: 9.0
Average Turnaround Time: 14.25

SJF:
Average Waiting Time: 9.0
Average Turnaround Time: 14.25

Priority:
Average Waiting Time: 9.0
Average Turnaround Time: 14.25

Round Robin:
Average Waiting Time: 11.25
Average Turnaround Time: 16.5

```

EXPERIMENT - 02

Aim: Write a program to simulate the following file allocation strategies.

(a) Sequential (b) Indexed (c) Linked

Algorithms:

1. Sequential File Allocation

- Search for Free Blocks: Start from the beginning of the file allocation table and look for consecutive free blocks.
- Allocate Blocks: If enough free blocks are found, mark them as allocated for the file.
- Update Tables: Store the file name and the corresponding blocks in the file allocation table and file list.

2. Indexed File Allocation

- Create Index Block: Allocate an index block for each file, which will hold the block numbers of the data blocks.
- Enter Block Numbers: For each data block, the user enters a block number.
- Allocate Blocks: Mark the entered block as allocated and store its number in the index block.
- Update Tables: Store the file name and the index block in the file allocation table and file list.

3. Linked File Allocation

- Enter Block Numbers: For each data block, the user enters a block number.
- Allocate Blocks: Mark the entered block as allocated and store the block numbers sequentially, creating a linked list of blocks.
- Update List: Store the file name and the list of block numbers in the file list. The file allocation table is not used in this strategy.

4. Displaying File Allocation Table

- Print Table: Show the file allocation table with block numbers and the corresponding file names.

5. Displaying Files

- Print Files: Show the list of files, including each file name and the blocks allocated to it.

Source Code:

```
class File: # Defines a class to represent a file with attributes name and size.
```

```

def __init__(self, name, size):
    self.name = name
    self.size = size

class Block: # Defines a class to represent a block with data and a reference
to the next block.
    def __init__(self, data):
        self.data = data
        self.next_block = None

def sequential_allocation(files, disk_blocks): # Allocates disk space to
files using sequential allocation.
    allocated_space = 0
    for file in files:
        if allocated_space + file.size <= disk_blocks:
            allocated_space += file.size
            print(f"Allocated {file.name} using sequential allocation")
        else:
            print(f"Not enough space to allocate {file.name} using sequential
allocation")

def indexed_allocation(files, disk_blocks): # Allocates disk space to files
using indexed allocation.
    index_blocks = len(files)
    if index_blocks <= disk_blocks:
        print("Allocating index blocks...")
        for file in files:
            print(f"Allocated index block for {file.name}")
            allocated_blocks = 0
            while allocated_blocks < file.size:
                print(f"Allocated data block for {file.name}")
                allocated_blocks += 1
    else:
        print("Not enough space to allocate index blocks for each file")

def linked_allocation(files, disk_blocks): # Allocates disk space to files
using linked allocation.
    print("Allocating data blocks...")
    for file in files:
        if file.size <= disk_blocks:
            allocated_blocks = 0
            current_block = Block(file.name)
            while allocated_blocks < file.size:
                print(f"Allocated data block for {file.name}")
                allocated_blocks += 1
                if allocated_blocks < file.size:
                    current_block.next_block = Block(file.name)
                    current_block = current_block.next_block

```

```

        else:
            print(f"Not enough space to allocate {file.name} using linked
allocation")

if __name__ == "__main__":
    disk_blocks = 10
    files = [
        File("readme.txt", 8),
        File("asset.txt", 4),
        File("design.txt", 5)
    ]
    print("Sequential Allocation:")
    sequential_allocation(files, disk_blocks)
    print("\nIndexed Allocation:")
    indexed_allocation(files, disk_blocks)
    print("\nLinked Allocation:")
    linked_allocation(files, disk_blocks)

```

Output:

```

PS C:\Users\D A GURUPRIYAN\Documents\MOS> & "C:/Users/D A GURUPRIYAN
Sequential Allocation:
Allocated readme.txt using sequential allocation
Not enough space to allocate asset.txt using sequential allocation
Not enough space to allocate design.txt using sequential allocation

Indexed Allocation:
Allocating index blocks...
Allocated index block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated index block for asset.txt
Allocated data block for asset.txt
Allocated data block for asset.txt
Allocated data block for asset.txt
Allocated data block for asset.txt
Allocated index block for design.txt
Allocated data block for design.txt
Allocated data block for design.txt
Allocated data block for design.txt
Allocated data block for design.txt
Allocated data block for design.txt

Linked Allocation:
Allocating data blocks...
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for readme.txt
Allocated data block for asset.txt
Allocated data block for asset.txt
Allocated data block for asset.txt
Allocated data block for asset.txt
Allocated data block for design.txt
Allocated data block for design.txt
Allocated data block for design.txt
Allocated data block for design.txt

```


EXPERIMENT - 03

Aim: Write a program to simulate paging technique of memory management

Algorithm:

- Start: Begin the process.
- Input Memory and Page Size: Ask the user to enter the memory size and page size in bytes.
- Calculate Number of Pages: Determine how many pages are needed by dividing the memory size by the page size and rounding up to the nearest whole number.
- Input Logical Address: Ask the user to enter a logical address.
- Calculate Page Offset: Find the page offset by taking the remainder of the logical address divided by the page size.
- Calculate Page Number: Determine the page number by dividing the logical address by the page size.
- Calculate Physical Address: Compute the physical address using the page number and page offset.
- Output Results: Display the number of pages, the page offset, and the physical address.
- Stop: End the process.

Source Code:

```
class PageTable: # Manages the page table for address translation.
    def __init__(self, num_frames): # Initializes the page table with a
specified number of frames.
        self.page_table = {}
        self.num_frames = num_frames

    def map_page(self, logical_address, frame_number): # Maps a logical
address to a frame number in the page table.
        self.page_table[logical_address] = frame_number
    def translate_address(self, logical_address): # Translates a logical
address to a physical address using the page table.
        if logical_address in self.page_table:
            return self.page_table[logical_address]
        else:
            return None

class Memory: # Represents the main memory.
    def __init__(self, num_frames): # Initializes the memory with a specified
number of frames.
        self.frames = [None] * num_frames
    def load_page(self, frame_number, page_data): # Loads a page into the
specified frame of memory.
        self.frames[frame_number] = page_data
```

```

    def get_page(self, frame_number): # Retrieves the data stored in a
specific frame of memory.
        return self.frames[frame_number]

class Process: # Represents a process with its associated pages.
    def __init__(self, process_id, pages): # Initializes the process with a
unique ID and its pages.
        self.process_id = process_id
        self.pages = pages

def simulate_paging(processes, memory_size): # Simulates paging by loading
processes into memory and performing address translation.
    num_frames = memory_size // 4 # Assuming each page is 4 units in size
    memory = Memory(num_frames)
    page_table = PageTable(num_frames)
    for process in processes:
        print(f"Loading process {process.process_id} into memory...")
        for i, page in enumerate(process.pages):
            frame_number = i % num_frames
            memory.load_page(frame_number, page)
            page_table.map_page((process.process_id, i), frame_number)
            print(f"Page {i} of process {process.process_id} loaded into frame
{frame_number}")
        print("\nAddress Translation:")
        for process in processes:
            print(f"Process {process.process_id} pages:")
            for i, _ in enumerate(process.pages):
                logical_address = (process.process_id, i)
                physical_address = page_table.translate_address(logical_address)
                print(f"Logical Address: {logical_address} -> Physical Address:
{physical_address}")

if __name__ == "__main__":
    processes = [
        Process(1, ['A', 'E', 'I', 'O', 'U']),
        Process(2, ['G', 'U', 'R', 'U', 'P']),
    ]
    memory_size = 16 # 16 units of memory
    simulate_paging(processes, memory_size)

```

Output:

```
● PS C:\Users\D A GURUPRIYAN\Documents\MOS> & "C:/Users/D A GU
py"
Loading process 1 into memory...
Page 0 of process 1 loaded into frame 0
Page 1 of process 1 loaded into frame 1
Page 2 of process 1 loaded into frame 2
Page 3 of process 1 loaded into frame 3
Page 4 of process 1 loaded into frame 0
Loading process 2 into memory...
Page 0 of process 2 loaded into frame 0
Page 1 of process 2 loaded into frame 1
Page 2 of process 2 loaded into frame 2
Page 3 of process 2 loaded into frame 3
Page 4 of process 2 loaded into frame 0

Address Translation:
Process 1 pages:
Logical Address: (1, 0) -> Physical Address: 0
Logical Address: (1, 1) -> Physical Address: 1
Logical Address: (1, 2) -> Physical Address: 2
Logical Address: (1, 3) -> Physical Address: 3
Logical Address: (1, 4) -> Physical Address: 0
Process 2 pages:
Logical Address: (2, 0) -> Physical Address: 0
Logical Address: (2, 1) -> Physical Address: 1
Logical Address: (2, 2) -> Physical Address: 2
Logical Address: (2, 3) -> Physical Address: 3
Logical Address: (2, 4) -> Physical Address: 0
```

EXPERIMENT - 04

Aim: Write a program to simulate the following file organization techniques

(a) Single level directory (b) Two level directory (c) Hierarchical.

Algorithm:

(a) Single Level Directory

1. Initialize: Start with an empty list to store file names and their sizes.
2. Operations:
 - create_file(file_name, size):
 - Check if the file name is already in the list. If it is, show an error message.
 - If not, add the file name and size to the list.
 - delete_file(file_name):
 - Check if the file name is in the list. If it is, remove it from the list.
 - If not, show an error message.
 - display_files():
 - Show all file names and their sizes from the list.

(b) Two Level Directory

1. Initialize: Start with an empty dictionary for the directory structure.
2. Operations:
 - create_file(directory_name, file_name, size):
 - Check if the directory exists. If not, create it as a key with an empty dictionary as its value.
 - Check if the file name is already in the directory. If it is, show an error message.
 - If not, add the file name and size to the directory.
 - delete_file(directory_name, file_name):
 - Check if the directory exists. If it does, check if the file name is in the directory. If it is, remove it.
 - If not, show an error message.
 - display_files():
 - Show the directory structure and files in it.

(c) Hierarchical Directory

1. Initialize: Start with a root directory as an empty dictionary.
2. Operations:
 - create_file(path, file_name, size):
 - Split the path into directory names.
 - Traverse the directory structure from the root using the directory names.
 - Create any directories that don't exist.
 - Check if the file name is already in the directory. If it is, show an error message.
 - If not, add the file name and size to the directory.
 - delete_file(path):
 - Split the path into directory names.
 - Traverse the directory structure from the root using the directory names.
 - Check if the file name is in the directory. If it is, remove it.
 - If not, show an error message.
 - display_files():
 - Show the directory structure and files in it.

Source Code:

```
class File: # Represents a file with attributes name and owner.
    def __init__(self, name, owner):
        self.name = name
        self.owner = owner

class SingleLevelDirectory: # Represents a single-level directory structure.
    def __init__(self): # Initializes an empty single-level directory.
        self.files = {}
    def add_file(self, file): # Adds a file to the single-level directory.
        self.files[file.name] = file
    def search_file(self, file_name): # Searches for a file in the single-
        level directory.
        if file_name in self.files:
            return self.files[file_name]
        else:
            return None

class TwoLevelDirectory: # Represents a two-level directory structure.
```

```

def __init__(self):
    self.master_directory = {}
def add_user_directory(self, user, directory):
    self.master_directory[user] = directory
def search_file(self, user, file_name):
    if user in self.master_directory:
        user_directory = self.master_directory[user]
        return user_directory.search_file(file_name)
    else:
        return None

class HierarchicalDirectory: # Represents a hierarchical directory structure.
    def __init__(self, name):
        self.name = name
        self.files = {}
        self.subdirectories = {}
    def add_file(self, file):
        self.files[file.name] = file
    def add_subdirectory(self, directory):
        self.subdirectories[directory.name] = directory
    def search_file(self, file_name):
        if file_name in self.files:
            return self.files[file_name]
        else:
            for subdirectory in self.subdirectories.values():
                result = subdirectory.search_file(file_name)
                if result:
                    return result
            return None

if __name__ == "__main__":
    print("Single Level Directory:")
    single_level_dir = SingleLevelDirectory()
    single_level_dir.add_file(File("single_level1.txt", "user1"))
    single_level_dir.add_file(File("single_level2.txt", "user2"))
    print(single_level_dir.search_file("single_level1.txt"))
    print("\nTwo Level Directory:")
    two_level_dir = TwoLevelDirectory()
    user1_dir = SingleLevelDirectory()
    user1_dir.add_file(File("two_level1.txt", "user1"))
    user2_dir = SingleLevelDirectory()
    user2_dir.add_file(File("two_level2.txt", "user2"))
    two_level_dir.add_user_directory("user1", user1_dir)
    two_level_dir.add_user_directory("user2", user2_dir)
    print(two_level_dir.search_file("user1", "two_level1.txt"))
    print("\nHierarchical Directory:")
    root_dir = HierarchicalDirectory("root")
    user1_dir = HierarchicalDirectory("user1")

```

```

user1_dir.add_file(File("file1.txt", "user1"))
user2_dir = HierarchicalDirectory("user2")
user2_dir.add_file(File("file2.txt", "user2"))
root_dir.add_subdirectory(user1_dir)
root_dir.add_subdirectory(user2_dir)
print(root_dir.search_file("file1.txt"))

```

Output:

```

● PS C:\Users\D A GURUPRIYAN\Documents\MOS> & "C:/U
n.py"
Single Level Directory:
<__main__.File object at 0x00000175BAC92BA0>

Two Level Directory:
<__main__.File object at 0x00000175BAC92FC0>

Hierarchical Directory:
<__main__.File object at 0x00000175BAC930B0>

```

EXPERIMENT - 05

Aim: Write a program to simulate Banker's algorithm for the purpose of deadlock avoidance.

Procedure:

- **Initialize Data Structures:**
 - **available_resources:** A list showing the number of available resources for each type.
 - **max_resources:** A table showing the maximum resources each process might need.
 - **allocated_resources:** A table showing the resources currently assigned to each process.
 - **need:** A table showing the remaining resources each process still needs.
 - **num_processes:** The number of processes.
 - **num_resources:** The number of resource types.
- **Calculate the Need Matrix:**
 - Subtract the **allocated_resources** from **max_resources** to get the **need** matrix.
- **Implement the is_safe_state Function:**
 - Copy **available_resources** to **work**.
 - Create a list called **finish** with all values set to False, indicating processes are not finished.
 - Create an empty list called **safe_sequence**.
 - While there are processes that can still get their required resources:
 - Check each process to see if its resource needs can be met by **work**.
 - If yes, update **work** with resources released by the process, mark it as finished, and add it to **safe_sequence**.
 - If all processes finish and **safe_sequence** includes all processes, return True and the sequence.
 - Otherwise, return False and an empty sequence.
- **Implement the request_resources Function:**
 - Check if the requested resources do not exceed the process's needs.
 - Check if the requested resources are available.

- If both checks pass, temporarily allocate the resources and update the state.
- Use `is_safe_state` to check if the system remains safe.
- If safe, print success and the sequence.
- If not safe, print a message indicating it leads to an unsafe state and undo the changes.
- **Implement the `display_state` Function:**
 - Print the current state of the system, showing the resources available, allocated, and needed for each process.

Source Code:

```
import numpy as np

class BankersAlgorithm: # Implements the Banker's Algorithm for deadlock avoidance.
    def __init__(self, allocation, max_demand, available):
        self.allocation = np.array(allocation)
        self.max_demand = np.array(max_demand)
        self.available = np.array(available)
        self.num_processes, self.num_resources = self.allocation.shape
    def is_safe_state(self): # Checks if the system is in a safe state.
        work = self.available.copy()
        finish = np.zeros(self.num_processes, dtype=bool)
        # Check if a process can be satisfied with the current resources
        def can_satisfy(process): # Checks if a process can be satisfied with the current resources.
            return all(need <= work) and not finish[process]
        safe_sequence = []
        while np.any(~finish):
            found = False
            for process in range(self.num_processes):
                need = self.max_demand[process] - self.allocation[process]
                if can_satisfy(process):
                    work += self.allocation[process]
                    finish[process] = True
                    safe_sequence.append(process)
                    found = True
                    break
            if not found:
                break
        return len(safe_sequence) == self.num_processes, safe_sequence

if __name__ == "__main__":
    allocation = [
        [1, 0, 0],
```

```

        [0, 2, 0],
        [2, 0, 3],
        [2, 3, 1],
        [2, 0, 2]
    ]
    max_demand = [
        [7, 5, 3],
        [3, 2, 2],
        [9, 0, 2],
        [2, 2, 2],
        [4, 3, 3]
    ]
    available = [2, 3, 2]
    banker = BankersAlgorithm(allocation, max_demand, available)
    safe, sequence = banker.is_safe_state()
    if safe:
        print("Safe Sequence:", sequence)
    else:
        print("Unsafe State, deadlock may occur")

```

Output:

```

PS C:\Users\D A GURUPRIYAN\Documents\MOS> &
Safe Sequence: [3, 1, 4, 0, 2]

```

EXPERIMENT - 06

Aim: Write a program to simulate disk scheduling algorithms

(a) FCFS (b) SCAN (c) C-SCAN.

Algorithm:

1. FCFS (First-Come, First-Served)

- Start at Initial Position: Begin from the initial head position.
- Process Requests in Order: Handle disk requests in the order they arrive.
- Calculate Movement: For each request, calculate the distance (absolute difference) between the current head position and the request.
- Update Head Position: Move the head to the position of the processed request.
- Repeat: Continue this for all requests.
- Total Movement: The total head movement is the sum of all calculated distances.

2. SCAN (Elevator) Algorithm

- Start at Initial Position: Begin from the initial head position.
- Move in One Direction: Move the head in one direction (towards higher or lower cylinder numbers).
- Process Requests: Handle requests as you encounter them, removing each from the list.
- Change Direction at End: When reaching the end of the disk, reverse the direction of the head.
- Repeat: Continue this process until all requests are handled.
- Total Movement: The total head movement is the total distance the head travels.

3. C-SCAN (Circular SCAN) Algorithm

- Start at Initial Position: Begin from the initial head position.
- Move in One Direction: Move the head in one direction to the end of the disk.
- Process Requests: Handle requests as you encounter them, removing each from the list.
- Jump to Beginning: After reaching the end, move the head back to the start of the disk.
- Continue in Same Direction: Move the head in the same direction again, processing any remaining requests.
- Repeat: Continue this process until all requests are handled.

- Total Movement: The total head movement is the total distance the head travels.

Source code:

```
class DiskScheduler: # Represents a disk scheduler for disk I/O operations.
    def __init__(self, requests, head_start):
        self.requests = requests
        self.head_position = head_start

    def fcfs(self): # Implements the First-Come, First-Served (FCFS) disk
scheduling algorithm.
        total_seek_time = 0
        for request in self.requests:
            total_seek_time += abs(request - self.head_position)
            self.head_position = request
        return total_seek_time

    def scan(self): # Implements the SCAN disk scheduling algorithm.
        total_seek_time = 0
        sorted_requests = sorted(self.requests)
        min_request = sorted_requests[0]
        max_request = sorted_requests[-1]
        # Move the head towards the maximum request
        for request in sorted_requests:
            if request >= self.head_position:
                total_seek_time += abs(request - self.head_position)
                self.head_position = request
        # Move the head towards the minimum request
        total_seek_time += abs(max_request - min_request)
        self.head_position = min_request
        return total_seek_time

    def c_scan(self): # Implements the Circular SCAN (C-SCAN) disk scheduling
algorithm.
        total_seek_time = 0
        sorted_requests = sorted(self.requests)
        max_request = sorted_requests[-1]
        # Move the head towards the maximum request
        for request in sorted_requests:
            if request >= self.head_position:
                total_seek_time += abs(request - self.head_position)
                self.head_position = request
        # Move the head to the beginning of the disk
        total_seek_time += abs(max_request)
        self.head_position = 0
        return total_seek_time

if __name__ == "__main__":
    requests = [120, 18, 37, 144, 14, 97, 65, 67]
```

```
head_start = 53
scheduler = DiskScheduler(requests, head_start)
print("FCFS Seek Time:", scheduler.fcfs())
print("SCAN Seek Time:", scheduler.scan())
print("C-SCAN Seek Time:", scheduler.c_scan())
```

Output:

```
● PS C:\Users\D A GURUPRIYAN\Documents\MOS>
y"
FCFS Seek Time: 542
SCAN Seek Time: 207
C-SCAN Seek Time: 274
```

EXPERIMENT - 07

Aim: Write a program to simulate page replacement algorithms

(a) FIFO (b) LRU (c) LFU

Procedure:

(a) FIFO (First-In, First-Out)

1. Initialize: Start with an empty page table.
2. Check Page Reference: For each page request, see if the page is in the page table.
3. Page Fault Handling:
 - If the page isn't in the table (page fault):
 - If there's space, add the page to the table.
 - If the table is full, remove the oldest page and add the new one.
4. Count Faults: Keep track of how many page faults occur.

(b) LRU (Least Recently Used)

1. Initialize: Start with an empty page table.
2. Check Page Reference: For each page request, see if the page is in the page table.
3. Page Fault Handling:
 - If the page isn't in the table (page fault):
 - If there's space, add the page to the table.
 - If the table is full, replace the least recently used page with the new one.
4. Update Usage: If the page is already in the table, mark it as the most recently used.
5. Count Faults: Keep track of how many page faults occur.

(c) LFU (Least Frequently Used)

1. Initialize: Start with an empty page table and a frequency table.
2. Check Page Reference: For each page request, see if the page is in the page table.
3. Page Fault Handling:
 - If the page isn't in the table (page fault):
 - If there's space, add the page to the table.
 - If the table is full, replace the least frequently used page with the new one.

4. Update Frequency: If the page is already in the table, increase its frequency count.
5. Count Faults: Keep track of how many page faults occur.

Source Code:

```
from collections import deque, Counter

class PageReplacement: # Represents a page replacement algorithm for managing
page faults.
    def __init__(self, page_frames):
        self.page_frames = page_frames
        self.page_queue = deque()
        self.page_access = []

    def fifo(self, page_references): # Implements the First-In, First-Out
(FIFO) page replacement algorithm.
        page_faults = 0
        for page in page_references:
            if page not in self.page_frames:
                page_faults += 1
                if len(self.page_frames) == len(self.page_queue):
                    if self.page_queue:
                        self.page_frames.remove(self.page_queue.popleft())
                self.page_frames.add(page)
                self.page_queue.append(page)
            elif page in self.page_frames:
                if page in self.page_queue:
                    self.page_queue.remove(page)
                self.page_queue.append(page)
        return page_faults

    def lru(self, page_references): # Implements the Least Recently Used (LRU)
page replacement algorithm.
        page_faults = 0
        for page in page_references:
            if page not in self.page_frames:
                page_faults += 1
                if len(self.page_frames) == len(self.page_access):
                    self.page_frames.remove(self.page_access.pop(0))
                self.page_frames.add(page)
                self.page_access.append(page)
            else:
                if page in self.page_access:
                    self.page_access.remove(page)
                self.page_access.append(page)
        return page_faults
```

```

def lfu(self, page_references): # Implements the Least Frequently Used
(LFU) page replacement algorithm.
    page_faults = 0
    page_counter = Counter()
    for page in page_references:
        if page not in self.page_frames:
            page_faults += 1
            if len(self.page_frames) == len(self.page_access):
                least_frequent_page = min(self.page_access,
key=page_counter.get)
                self.page_frames.remove(least_frequent_page)
                self.page_access.remove(least_frequent_page)
                self.page_frames.add(page)
                self.page_access.append(page)
            page_counter[page] += 1
    return page_faults

if __name__ == "__main__":
    page_references = [2, 2, 3, 4, 5, 1, 2, 4, 9, 7, 8, 7, 8, 1, 7, 8, 9, 4,
4, 5, 4, 2, 3, 2, 2, 2, 3, 6]
    page_frames = set()
    pr = PageReplacement(page_frames)
    print("FIFO Page Faults:", pr.fifo(page_references.copy()))
    pr = PageReplacement(page_frames)
    print("LRU Page Faults:", pr.lru(page_references.copy()))
    pr = PageReplacement(page_frames)
    print("LFU Page Faults:", pr.lfu(page_references.copy()))

```

Output:

```

● PS C:\Users\D A GURUPRIYAN\Documents\MOS>
algo.py"
FIFO Page Faults: 24
LRU Page Faults: 8
LFU Page Faults: 0

```


EXPERIMENT - 08

Aim: Write a program to simulate producer-consumer problem using semaphores

Procedure:

- **Initialization:**
 - Define buffer size and initialize buffer, mutex semaphore, empty semaphore (equal to buffer size), and full semaphore (set to 0).
 - Define the number of items to produce and initialize counters for items produced and consumed.
- **Producer Class:**
 - Create a class called Producer extending threading.Thread.
 - Implement the run() method for the Producer:
 - Loop until desired items are produced.
 - Simulate production time with sleep function.
 - Acquire empty semaphore to wait if buffer is full.
 - Acquire mutex semaphore to access buffer.
 - Append produced item to buffer.
 - Print producer's thread name and the produced item.
 - Update items produced counter.
 - Release mutex semaphore.
 - Release full semaphore to notify consumer an item is available.
- **Consumer Class:**
 - Create a class called Consumer extending threading.Thread.
 - Implement the run() method for the Consumer:
 - Loop until desired items are consumed.
 - Simulate consumption time with sleep function.
 - Acquire full semaphore to wait if buffer is empty.
 - Acquire mutex semaphore to access buffer.
 - Remove consumed item from buffer.
 - Print consumer's thread name and the consumed item.
 - Update items consumed counter.
 - Release mutex semaphore.

- Release empty semaphore to notify producer that space is available.
- **Create Instances:**
 - Create instances of Producer and Consumer classes.
- **Start Threads:**
 - Start producer and consumer threads.
- **Wait for Completion:**
 - Wait for all threads to finish using join() method on each thread.

Source Code:

```
import time
import random

BUFFER_SIZE = 5

buffer = []
buffer_lock = False
buffer_full = 0

def producer_consumer(): # Defines the producer-consumer function.
    global buffer, buffer_lock, buffer_full
    while True:
        item = random.randint(1, 100)
        if buffer_full == BUFFER_SIZE:
            print("Buffer is full. Waiting...")
            time.sleep(0.1) # Wait if buffer is full
        else:
            if buffer_lock:
                print("Buffer is being modified. Waiting...")
                time.sleep(0.1) # Wait if buffer is being modified
            else:
                buffer_lock = True # Acquire the lock
                buffer.append(item) # Add item to the buffer
                buffer_full += 1
                print(f"Produced {item}, Buffer: {buffer}")
                buffer_lock = False # Release the lock
                if random.random() < 0.5:
                    if buffer_full == 0:
                        print("Buffer is empty. Waiting...")
                        time.sleep(0.1) # Wait if buffer is empty
                    else:
                        if buffer_lock:
                            print("Buffer is being modified. Waiting...")
```

```

        time.sleep(0.1) # Wait if buffer is being
modified
        else:
            buffer_lock = True # Acquire the lock
            item = buffer.pop(0) # Remove item from the
buffer

            buffer_full -= 1
            print(f"Consumed {item}, Buffer: {buffer}")
            buffer_lock = False # Release the lock

        time.sleep(random.uniform(0.1, 1))
# Start the producer-consumer loop
producer_consumer()

```

Output:

```

⊗ PS C:\Users\D A GURUPRIYAN\Documents\MOS> & "C:/Users/D A G
.py"
Produced 78, Buffer: [78]
Produced 58, Buffer: [78, 58]
Consumed 78, Buffer: [58]
Produced 59, Buffer: [58, 59]
Consumed 58, Buffer: [59]
Produced 100, Buffer: [59, 100]
Produced 46, Buffer: [59, 100, 46]
Consumed 59, Buffer: [100, 46]
Produced 98, Buffer: [100, 46, 98]
Consumed 100, Buffer: [46, 98]
Produced 32, Buffer: [46, 98, 32]
Produced 57, Buffer: [46, 98, 32, 57]
Consumed 46, Buffer: [98, 32, 57]
Produced 75, Buffer: [98, 32, 57, 75]
Consumed 98, Buffer: [32, 57, 75]
Produced 73, Buffer: [32, 57, 75, 73]
Produced 36, Buffer: [32, 57, 75, 73, 36]
Buffer is full. Waiting...

```

EXPERIMENT - 09

Aim: Write a program to simulate the concept of Dining-Philosophers problem

Procedure:

1. Setup: Create a group of n philosophers and n chopsticks.
2. Initialization:
 - Each philosopher gets a unique identifier and references to their left and right chopsticks.
3. Start Threads:
 - Begin a thread for each philosopher.
4. Philosopher's Lifecycle:
 - Each philosopher enters an infinite loop representing their lifecycle:
 1. Think: Output that the philosopher is thinking.
 2. Wait for Chopsticks: Wait until both chopsticks are available.
 3. Pick Up Chopsticks: Acquire locks for both the left and right chopsticks.
 4. Eat: Output that the philosopher is eating.
 5. Put Down Chopsticks: Release locks for both the left and right chopsticks.
5. Repeat Indefinitely:
 - Repeat this loop indefinitely.

Source code:

```
import time
import random

NUM_PHILOSOPHERS = 5
MAX_EATING_TIME = 3

forks = [True] * NUM_PHILOSOPHERS # True indicates the fork is available

def philosopher(index): # Defines the philosopher function.
    for _ in range(MAX_EATING_TIME):
        think(index)
        eat(index)

def think(index): # Defines the think function.
    print(f"Philosopher {index} is thinking...")
    time.sleep(random.uniform(0.1, 1))
```

```

def eat(index): # Defines the eat function.
    print(f"Philosopher {index} is hungry and trying to pick up forks...")
    left_fork = index
    right_fork = (index + 1) % NUM_PHILOSOPHERS
    while True:
        if forks[left_fork] and forks[right_fork]:
            forks[left_fork] = False
            forks[right_fork] = False
            print(f"Philosopher {index} is eating...")
            time.sleep(random.uniform(0.1, 1))
            forks[left_fork] = True
            forks[right_fork] = True
            break
        else:
            print(f"Philosopher {index} couldn't acquire forks. Retrying...")
            time.sleep(0.1)

if __name__ == "__main__":
    for i in range(NUM_PHILOSOPHERS):
        philosopher(i)
    print("Dinner is over. Philosophers are full and content.")

```

Output:

```

PS C:\Users\D A GURUPRIYAN\Documents\MOS> & "C:/Users/D A GURUPRIYAN/App
Philosopher 0 is thinking...
Philosopher 0 is hungry and trying to pick up forks...
Philosopher 0 is eating...
Philosopher 0 is thinking...
Philosopher 0 is hungry and trying to pick up forks...
Philosopher 0 is eating...
Philosopher 0 is thinking...
Philosopher 0 is hungry and trying to pick up forks...
Philosopher 0 is eating...
Philosopher 1 is thinking...
Philosopher 1 is hungry and trying to pick up forks...
Philosopher 1 is eating...
Philosopher 1 is thinking...
Philosopher 1 is hungry and trying to pick up forks...
Philosopher 1 is eating...
Philosopher 1 is thinking...
Philosopher 1 is hungry and trying to pick up forks...
Philosopher 1 is eating...
Philosopher 2 is thinking...
Philosopher 2 is hungry and trying to pick up forks...
Philosopher 2 is eating...
Philosopher 2 is thinking...
Philosopher 2 is hungry and trying to pick up forks...
Philosopher 2 is eating...
Philosopher 2 is thinking...
Philosopher 2 is hungry and trying to pick up forks...
Philosopher 2 is eating...
Philosopher 3 is thinking...
Philosopher 3 is hungry and trying to pick up forks...
Philosopher 3 is eating...
Philosopher 3 is thinking...
Philosopher 3 is hungry and trying to pick up forks...
Philosopher 3 is eating...
Philosopher 3 is thinking...
Philosopher 3 is hungry and trying to pick up forks...
Philosopher 3 is eating...
Philosopher 4 is thinking...
Philosopher 4 is hungry and trying to pick up forks...
Philosopher 4 is eating...
Philosopher 4 is thinking...
Philosopher 4 is hungry and trying to pick up forks...
Philosopher 4 is eating...
Philosopher 4 is thinking...
Philosopher 4 is hungry and trying to pick up forks...
Philosopher 4 is eating...
Dinner is over. Philosophers are full and content.

```